

CRYPTOAGENT · A PLAIN READING

The Trick

Everything this machine does, explained the way I'd explain a card trick: slowly, and with the mystery removed. There isn't any mystery. That's the point.

If a thing can't be explained in plain words, it usually can't be explained at all — someone is just hoping you won't ask.

Written for Tanish · a guide to what the terminal is actually doing

What you're actually looking at

Let's start honestly. A memecoin is not a company. It has no revenue, no product, and no reason to be worth anything at all. What it has is a story and an audience.

So stop asking "is this a good asset." It isn't. Nothing here is. The only question that has ever mattered in this market is a different one:

What does the crowd currently want to believe — and how much of that belief has already been spent?

That's a confidence game. I don't mean that as an insult; I mean it as a description of the mechanism. Somebody builds a story, an audience arrives because the story is fun and the chart is moving, and the people who arrived early sell to the people who arrived late. Everyone involved knows this. They're each betting they're early.

Once you see it that way, the machine gets simple. It isn't valuing anything. It's watching a crowd, and trying to work out where in the performance we are.

The five things it watches


LAYER	THE QUESTION IT ANSWERS
Macro	Is there any money in the room at all? Solana memecoins are a leveraged bet on Solana. When the chain bleeds, every clever chart below it is a trap.
Structure	Is the pool deep enough that a move means something, or is one buyer pushing a puddle around?
Crowd	Is attention arriving, and is there any left to arrive?
Social	Who is saying it, and have they ever been right before?
Forecast	Given all that — what do we actually claim, out loud, before we find out?

The rule everything else hangs from. If a number on that screen can't be traced back to something observable, it isn't insight. It's decoration. Every score the machine produces carries the chain of observations that made it. When the chain is empty, it says so and declines to guess. A confident number with no reasoning behind it is exactly how people get fooled — usually by themselves.

Attention is a hill, not a ladder

Here's the mistake almost every one of these systems makes, and it makes it confidently: it treats attention as a straight line. More people talking equals better. More mentions, higher score.

Think about that for a second and you'll see the problem. If everybody already knows, who exactly is left to buy it from you?



Attention, measured against what's normal for a coin that size. The machine puts the peak at roughly three times the typical level for its tier.

Too little attention and there's no bid — you've found something genuinely early and genuinely stuck. Too much and the price already contains everyone's enthusiasm; you'd be buying the applause at the end of the show.

A system that scores attention as a straight line will, with total confidence and perfect consistency, buy tops. Not occasionally. By design.

Something I got wrong, and it's instructive. The first version of this curve flattened everything above three times normal, so a coin at *thirty-nine* times normal scored almost the same as one at three. Wildly saturated and perfectly placed looked identical. The fix was arithmetic, not insight — but nobody would have noticed, because the number looked reasonable. Numbers usually do.


Where people actually buy

Ask someone when the buying pressure peaks and they'll say the top. It doesn't. At the top, most people are watching, not acting. They feel late. Feeling late makes people hesitate.

The heaviest buying comes slightly afterwards, and the reason is entirely emotional:

Somebody watched it run without them. Then it dipped. Now it's turning back up — and this time they're not going to miss it.

Regret, plus a discount. That combination moves money faster than any chart pattern, because it isn't a chart pattern. It's a person who has already decided, waiting for permission.



Ran hard, paused, turning up again. The machine looks for all three at once — any one of them alone means nothing.


So the machine looks for that exact shape: a real run over six hours, a cooling or flat stretch over one, and a tick upward in the last five minutes. All three together. A rise alone is just a rise. A dip alone is just a dip. It's the sequence that carries the meaning, because the sequence is what a person experiences.

Ignition and fuel

Two questions that sound like one question, and aren't:

- Will the crowd arrive in the next few minutes? Call that *FOMO*. Ignition.
- Will this still be burning in two or three hours? Call that *Aura*. Fuel.

Keep them separate, because they come apart constantly — and the way they come apart is the useful part.



The top-right corner is where a system that only measures excitement spends all of its money.

Rank on excitement alone and you will systematically buy the top-left box: the loud, hollow ones. They look wonderful right up until they don't.


From the live terminal, the night this was written. Four coins scored high on durability. Every single one of them had almost no crowd — one read **88 for fuel and 16 for ignition**. Nobody was talking about them at all. A system watching only for excitement would never have shown you a single one.

The clock, not the label

This one is yours. You watched a coin die and told me the shape of it, and you were right — and being right exposed something the machine was doing badly.

It had been treating rug risk as a *property*. A thing you check once: is the pool locked, can they mint more, yes or no, filed away for half an hour. That's the wrong shape entirely. Rug risk isn't a property.

It's a countdown. And the tell isn't in the contract or the price — it's in the liquidity.



The bait is other people's money. The clock starts when attention does.

Somebody has to be attracted into the pool before there's anything worth stealing. So the mention isn't a buy signal at all — it's the start of a timer. The danger rises as the pool fills, peaks in a band where it's fat enough to be worth burning a wallet for but not so fat that anyone's watching, and then it ends.

The machine now tracks four things and a gate. The gate is **capability** — if the pool is genuinely locked and the authorities are renounced, nobody can pull it and the clock doesn't run. Then: is the pool in the harvest band, is liquidity flowing in right now, and is the price outrunning the pool so buyers are crowding thin ground.

Now the part I have to say plainly. In this machine's own records, liquidity collapse happened in about **three percent** of graded cases. A detector for a three-percent event will be wrong most times it opens its mouth. That's not pessimism, it's arithmetic.


So the clock writes down every warning it makes and gets marked against reality like everything else. If it can't beat simply assuming nothing bad happens, it should be switched off — and the record will say so out loud. A hunch that refuses to be tested is just a hunch with good posture.

Nothing gets to be an opinion

Here's the part that matters more than any clever signal, and it's the least exciting part to look at.

Anyone can sound confident about a coin. Confidence is free. So nothing in this machine is allowed to simply *have a view*. Every component — every model, every Telegram channel, even the plain arithmetic score — has to say something specific enough to be wrong:

- Which direction.
- By when, exactly.
- How sure, as a number.
- What would prove it wrong.



Not a clever loop. Just the only honest one.

Then the machine goes back on its own, later, and checks. It doesn't wait to be asked and it doesn't wait for a good mood. That record is the single thing here that cannot be regenerated if it's lost, which is why it's the one file worth protecting.

The plain benchmark. The simple arithmetic score competes in the same table as everything expensive. If a sophisticated model can't beat plain arithmetic, the honest response is to turn the model off — not to explain why the comparison is unfair.

And a caution about early numbers. Under fifty or so graded calls, a track record is noise wearing a percentage sign. The table says *unproven* rather than showing you a flattering figure. It'll go on saying that for days. Let it.

When somebody is right too often

Suppose one of your channels starts calling them correctly. Really correctly, more than luck would allow. The instinct is obvious: follow it harder.

Resist that, because there are only two ordinary explanations and they point in opposite directions.

They read the tape

The signals were already visible when they spoke.

*Worth following —
and worth copying.*

They knew first

The tape was flat and it moved anyway.

*You are not a follower.
You are the exit.*

Same excellent hit rate. Completely opposite meaning.

Telling them apart doesn't require anything exotic, which is usually the case with things that look like mind-reading. The machine already recorded what the market looked like at the moment each call was made. So it asks a simple question of the winners: *was anything visible yet?*

If volume was already climbing and the price was already moving, they were reading the same screen you can read. That's a skill, it's followable, and you could learn it. If everything was flat and it moved anyway — they weren't reading. They were involved.

There's a third case worth knowing, particular to this market: a large enough channel can *cause* the move it predicted. That's neither skill nor secret knowledge. It's just a crowd being pointed at something.

Three questions, not one

The thing most likely to jump in eight minutes is rarely the thing most likely to survive the week. Rank them together and you've answered neither question. So the board asks three separate ones, and admits they're not equally answerable.

WINDOW	WHAT IT'S REALLY ASKING	HOW MUCH TO TRUST IT
5–10 minutes	Is it igniting right now — flow, acceleration, crowd	Fast and shallow. Fine, if you're watching.
1–5 hours	Does the trend have fuel — structure, absorption, participation	The strongest thing here.
1–10 days	Will it still exist	A screen, not a forecast.

That last row is the honest one. There is no tested long-range signal in this machine. Its longest measured window is six hours, and most of these coins will not exist in ten days. So it doesn't pretend to forecast returns out there — it ranks the few traits that correlate with still being tradeable later, and it labels itself a screen everywhere it appears.

Three tiers presented as equally solid would be a lie told with good formatting. Two of them are readings. One is a filter.

You'll also notice a tier occasionally says *nothing clears the bar* and names the best thing it rejected. That's deliberate. Filling a slot because a slot exists is how a system starts lying to you politely.

What it can't do

- **It can't see who holds the coin.** Wallet concentration, insider clusters, connected wallets — all invisible without paid data. That's the largest blind spot.
- **It can't read a private channel.** Only public ones, and the public ones you have wired mostly post news, not calls.
- **It can't tell you a thing is safe.** Nothing here is safe. It can sometimes tell you a thing is *more dangerous than it looks*, which is a different and more useful service.
- **It cannot buy anything.** No keys, no wallet, no orders. It produces a reading and a record. What you do next is entirely yours.

Last thought, and it's the only advice in here. The machine's real value isn't any single number on the screen — it's that it writes down what it thought *before* it knew, and then goes back and checks. Do that yourself and you'll improve faster than any model I could build you. Most people never write it down, which is precisely why most people remember themselves as having been right.